

▼ Nhu M Vo

Lab #2

```
from __future__ import division
import pandas as pd
import numpy as np
import statsmodels.api as sm
import statsmodels.formula.api as smf
import os
import matplotlib.pyplot as plt

from google.colab import files
uploaded = files.upload()

import io
d = pd.read_csv(io.BytesIO(uploaded['GSS.2006.csv']))

d.head()
```

[Choose Files](#) GSS.2006.csv

- **GSS.2006.csv**(text/csv) - 8232424 bytes, last modified: 6/7/2023 - 100% done
Saving GSS.2006.csv to GSS.2006.csv

| | vpsu | vstrat | adults | ballot | dateintv | famgen | form | formwt | gender1 | hompop | ... | away7 | gender14 | old14 | rela |
|---|------|--------|--------|--------|----------|--------|------|--------|---------|--------|-----|-------|----------|-------|------|
| 0 | 1 | 1957 | 1 | 3 | 316 | 2 | 1 | 1 | 2 | 3 | ... | NaN | NaN | NaN | |
| 1 | 1 | 1957 | 2 | 2 | 630 | 1 | 2 | 1 | 2 | 2 | ... | NaN | NaN | NaN | |
| 2 | 1 | 1957 | 2 | 2 | 314 | 2 | 1 | 1 | 2 | 2 | ... | NaN | NaN | NaN | |
| 3 | 1 | 1957 | 1 | 1 | 313 | 1 | 2 | 1 | 2 | 1 | ... | NaN | NaN | NaN | |
| 4 | 1 | 1957 | 3 | 1 | 322 | 2 | 2 | 1 | 2 | 3 | ... | NaN | NaN | NaN | |

5 rows × 1261 columns



1. Recode 2 variable into new categories. They can both be continuous-ish or both be nominal-ish, or one of each. Tell me what you did and explain the variable(s).
- ▼
- ▼ Create a number of categories for frequency of going to a bar/tavern: low, medium, high

```
## 3 options of low, medium, high ##

conditions = [
    (d['socbar'] < 3) & (d['socbar'] > 0),
    (d['socbar'] > 2) & (d['socbar'] < 6),
    (d['socbar'] > 5)]
choices = [1, 2, 3]
d['cut'] = np.select(conditions, choices, default=np.nan)

# Look at the results

d.cut.describe()
```

```

count      1989.000000
mean        2.582705
std         0.609138
min         1.000000
25%         2.000000
50%         3.000000
75%         3.000000
max         3.000000
Name: cut, dtype: float64

```

```
## How many of each category are there?
```

```

d.cut.value_counts()

3.0      1286
2.0       576
1.0       127
Name: cut, dtype: int64

```

- Check the recoding: It is "1" for categories more than 0 and less than 3, it is "2" for categories from 3 to 5, and it is "3" for categories greater than 5 (high, medium, low, respectively). The columns for each 1,2,3 add up to 100%

Note: 1 - almost everyday, 2 - once of twice a week, 3 - several times a month, 4 - about once a month, 5 - several times a year, 6 - about once a year

```

res = pd.crosstab(d.socbar, d.cut)
res.astype('float').div(res.sum(axis=0), axis=1)

```

| | cut | 1.0 | 2.0 | 3.0 |
|--------|-----|----------|----------|--------|
| socbar | | | | |
| 1.0 | | 0.110236 | 0.000000 | 0.0000 |
| 2.0 | | 0.889764 | 0.000000 | 0.0000 |
| 3.0 | | 0.000000 | 0.211806 | 0.0000 |
| 4.0 | | 0.000000 | 0.347222 | 0.0000 |
| 5.0 | | 0.000000 | 0.440972 | 0.0000 |
| 6.0 | | 0.000000 | 0.000000 | 0.2014 |
| 7.0 | | 0.000000 | 0.000000 | 0.7986 |

#Simply another way to do the same thing as above: A nice function someone wrote to do the same thing:

```

def binning(col, cut_points, labels=None):
    #Define min and max values:
    minval = col.min()
    maxval = col.max()

    #create list by adding min and max to cut_points
    break_points = [minval] + cut_points + [maxval]

    #if no labels provided, use default labels 0 ... (n-1)
    if not labels:
        labels = range(len(cut_points)+1)

    #Binning using cut function of pandas
    colBin = pd.cut(col,bins=break_points,labels=labels,include_lowest=True)
    return colBin

```

```
#Binning attend:

cut_points = [2,5]
labels = ["low","medium","high"]
d["socbar_cut"] = binning(d["socbar"], cut_points, labels)
##print pd.value_counts(d["attend_cut"], sort=False)##
```

```
## See it works the same way...
```

```
summary = d.socbar_cut.describe()
summary = summary.transpose()
summary

count      1989
unique        3
top         high
freq       1286
Name: socbar_cut, dtype: object
```

```
## See it works the same way...
```

```
d.socbar_cut.value_counts()

high      1286
medium     576
low        127
Name: socbar_cut, dtype: int64
```

```
## See it works the same way...
```

```
res = pd.crosstab(d.socbar, d.socbar_cut)
res.astype('float').div(res.sum(axis=0), axis=1)
```

| | socbar_cut | low | medium | high |
|--------|------------|----------|----------|--------|
| socbar | | | | |
| 1.0 | | 0.110236 | 0.000000 | 0.0000 |
| 2.0 | | 0.889764 | 0.000000 | 0.0000 |
| 3.0 | | 0.000000 | 0.211806 | 0.0000 |
| 4.0 | | 0.000000 | 0.347222 | 0.0000 |
| 5.0 | | 0.000000 | 0.440972 | 0.0000 |
| 6.0 | | 0.000000 | 0.000000 | 0.2014 |
| 7.0 | | 0.000000 | 0.000000 | 0.7986 |

- Below is a binary recode for "how often you go to a bar/tavern": with 0 being <5 times, and 1 being more than 4 times OR everything greater than or equal to 5

```
##Simple binary cut##

conditions = [
    (d['socbar'] < 5) ,
    (d['socbar'] >= 4 ) ]
choices = [0,1]
d['high'] = np.select(conditions, choices, default=np.nan)
```

- Check the recoding: It is 0 for all categories less than 5 and 1 for everything greater than or equal to 5

```
res = pd.crosstab(d.socbar, d.high)
res.astype('float').div(res.sum(axis=0), axis=1)
```

| | high | 0.0 | 1.0 |
|--------|----------|----------|-----|
| socbar | | | |
| 1.0 | 0.031180 | 0.000000 | |
| 2.0 | 0.251670 | 0.000000 | |
| 3.0 | 0.271715 | 0.000000 | |
| 4.0 | 0.445434 | 0.000000 | |
| 5.0 | 0.000000 | 0.164935 | |
| 6.0 | 0.000000 | 0.168182 | |
| 7.0 | 0.000000 | 0.666883 | |

- Let's look at another variable - this is a question about: "Have you ever had sex with someone other than your husband or wife while you were married?"(evstray) with 1 being Yes, 2 being No. The second variable in this Lab Assignment that I'm going to recode is getting rid of the "Never married" (column 3).

```
## Table 2 ##
```

```
my_tab = pd.crosstab(index=d["evstray"], # Make a crosstab
                     columns="count")
```

```
d.loc[d['evstray'] == 3.0, 'evstray'] = np.nan
#look into location where value in evstray=3, replace it with np.nan (NaN value)
```

```
display(my_tab)
def compute_percentage(x):
    pct = float(x/my_tab['count'].sum()) * 100
    return round(pct, 2)
```

```
my_tab['percentage'] = my_tab.apply(compute_percentage, axis=1)
```


```
my_tab
```

| col_0 | count | |
|---------|-------|------------|
| evstray | | |
| 1.0 | 350 | |
| 2.0 | 1414 | |
| 3.0 | 623 | |
| col_0 | count | percentage |
| evstray | | |
| 1.0 | 350 | 14.66 |
| 2.0 | 1414 | 59.24 |
| 3.0 | 623 | 26.10 |

```
d.hist(column='evstray')
```



```
array([[<Axes: title={ 'center': 'evstray' }>]], dtype=object)
```



2. Use one (or both) of your recoded variables to do a cross-tabulation. Explain your results.

"evstray" (Independent variable: whether or not one cheats, with 1 being yes, and 2 being no) is the column variable and "cut" (Dependent variable: how often an individual goes to a bar/tavern, with 1 being high frequency and 3 being low frequency) is the row variable. The below results indicate that for the people who cheat on their spouse, 9.09% are going to the bar frequently while for non-cheaters, 4.13% of them are going to the bar frequently, which is contrary with my hypothesis because people cheat on their spouses at the bar/tavern.

```
res = pd.crosstab(d.cut, d.evstray)
res.astype('float').div(res.sum(axis=0), axis=1)
```

```
display(d.evstray)
```

```
0      NaN
1      NaN
2      NaN
3      NaN
4      NaN
...
4505    2.0
4506    2.0
4507    NaN
4508    NaN
4509    2.0
Name: evstray, Length: 4510, dtype: float64
```

```
res = pd.crosstab(d.cut, d.evstray)
res.astype('float').div(res.sum(axis=0), axis=1)
```

| evstray | 1.0 | 2.0 |
|---------|----------|----------|
| cut | | |
| 1.0 | 0.090909 | 0.041322 |
| 2.0 | 0.367965 | 0.253099 |
| 3.0 | 0.541126 | 0.705579 |

- Below, I'm using pandas "crosstab" function to get column percentages now.

(I put the addition "*100" at the end because I wanted it to show up as percentages)

```
res.p = pd.crosstab(d['cut'], d['evstray'], normalize='columns')*100
res.p
```

```
<ipython-input-18-dfeb775cf6e9>:1: UserWarning: Pandas doesn't allow columns to be created via a new attribute
res.p = pd.crosstab(d['cut'], d['evstray'], normalize='columns')*100
```

| evstray | 1.0 | 2.0 |
|---------|-----------|-----------|
| cut | | |
| 1.0 | 9.090909 | 4.132231 |
| 2.0 | 36.796537 | 25.309917 |
| 3.0 | 54.112554 | 70.557851 |

3. Run a linear regression with 1 independent and 1 dependent variable; make all of

- the recodes necessary to make the model as easy to interpret as possible; and explain your results.

Asking the question: How do you think "evstray" - whether or not one cheats on their spouse (IV) should be

- related to "cut" - their frequency of going to a bar/tavern (DV). In other words, how does the likelihood of cheating on a spouse (Yes/No) affect their frequency of going to the bar/tavern

```
lm = smf.ols(formula = 'evstray~cut', data = d).fit()
print (lm.summary())
```

| OLS Regression Results | | | | | | |
|------------------------|------------------|---------------------|--------|-----------|--------|--------|
| ===== | | | | | | |
| Dep. Variable: | evstray | R-squared: | | 0.021 | | |
| Model: | OLS | Adj. R-squared: | | 0.020 | | |
| Method: | Least Squares | F-statistic: | | 25.89 | | |
| Date: | Mon, 12 Jun 2023 | Prob (F-statistic): | | 4.20e-07 | | |
| Time: | 02:08:32 | Log-Likelihood: | | -572.91 | | |
| No. Observations: | 1199 | AIC: | | 1150. | | |
| Df Residuals: | 1197 | BIC: | | 1160. | | |
| Df Model: | 1 | | | | | |
| Covariance Type: | nonrobust | | | | | |
| ===== | | | | | | |
| | coef | std err | t | P> t | [0.025 | 0.975] |
| ----- | | | | | | |
| Intercept | 1.5479 | 0.052 | 29.642 | 0.000 | 1.445 | 1.650 |
| cut | 0.0989 | 0.019 | 5.088 | 0.000 | 0.061 | 0.137 |
| ===== | | | | | | |
| Omnibus: | 264.735 | Durbin-Watson: | | 1.996 | | |
| Prob(Omnibus): | 0.000 | Jarque-Bera (JB): | | 468.245 | | |
| Skew: | -1.515 | Prob(JB): | | 2.10e-102 | | |
| Kurtosis: | 3.441 | Cond. No. | | 14.1 | | |
| ===== | | | | | | |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Conclusion: Knowing whether or not one cheats do not help predict how frequent they go to the bar/tavern, and vice versa. A coefficient of 0.09 indicates no obvious relationship between whether one cheats (Yes/No)

and frequency of going to a bar/tavern. The p-values are - which means the results are statistically significant.

4. Plot two variables, either as a scatter plot or boxplot; add in trend/regression lines; and explain your results.

Here is plotting of "evstray" (whether or not one cheats) against "cut" (how often an individual goes to bar/tavern).

Note:


"cut": (Dependent variable: how often an individual goes to a bar/tavern, with 1 being high frequency and 3 being low frequency)

"evstray": (Independent variable: whether or not one cheats, with 1 being yes, and 2 being no)

The results below show that the individuals that cheat on their spouses go to tavern/bars less frequently than their counterparts who don't cheat


```
import seaborn as sns
sns.regplot(x=d['evstray'], y=d['cut'])
```

<Axes: xlabel='evstray', ylabel='cut'>



```
sns.jointplot(data=d, x='evstray', y='cut', kind="reg")
```

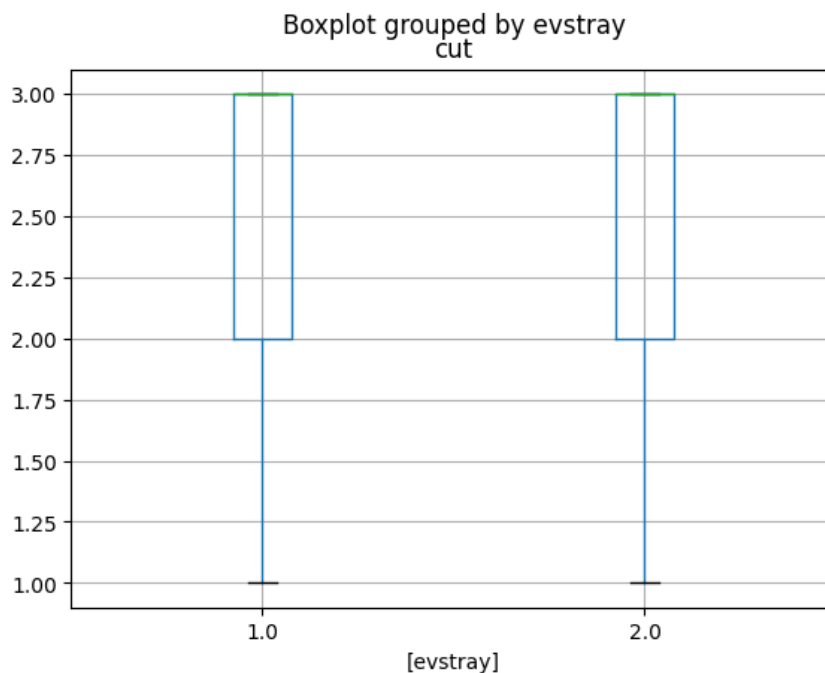
```
<seaborn.axisgrid.JointGrid at 0x7f3ff35cf010>
```



this jointplot with a negative regression line shows that the relationship between whether or not one cheats on their spouse is negatively correlated with how often they go to a pub/tavern

```
1.00 1 2
%matplotlib inline
d.boxplot(column='cut', by=['evstray'])

<Axes: title={'center': 'cut'}, xlabel='[evstray]'
```



✓ 1s completed at 10:08 PM

● ×